

```
# -----
# 1. IMPORT LIBRARIES
# -----
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from tensorflow.keras.datasets import fashion_mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.metrics import classification_report, confusion_matrix
import tensorflow as tf
```

```
# Load Fashion MNIST dataset
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()


class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
               'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

print("Training samples:", x_train.shape)
print("Test samples:", x_test.shape)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-dataset-29515/29515 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-dataset-26421880/26421880 2s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-dataset-5148/5148 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-dataset-4422102/4422102 1s 0us/step
Training samples: (60000, 28, 28)
Test samples: (10000, 28, 28)
```

Visualize Sample images

```
plt.figure(figsize=(10,5))
for i in range(10):
    plt.subplot(2,5,i+1)
    plt.imshow(x_train[i], cmap='gray')
    plt.title(class_names[y_train[i]])
    plt.axis('off')
plt.show()
```



```
# Normalize images
x_train = x_train / 255.0
x_test = x_test / 255.0

# One-hot encode labels
y_train_cat = to_categorical(y_train, 10)
y_test_cat = to_categorical(y_test, 10)

# OPTIONAL: Data augmentation
datagen = ImageDataGenerator(
    rotation_range=10,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=False
)
datagen.fit(x_train.reshape(-1, 28, 28, 1))
```

BUILD FEEDFORWARD NEURAL NETWORK

```
model = Sequential([
    Flatten(input_shape=(28,28)),
    Dense(256, activation='relu'),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])
```

```

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

model.summary()

```

```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/reshaping/flatten.py:37
    super().__init__(**kwargs)
Model: "sequential"

```

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 256)	200,960
dense_1 (Dense)	(None, 128)	32,896
dense_2 (Dense)	(None, 10)	1,290

```

Total params: 235,146 (918.54 KB)
Trainable params: 235,146 (918.54 KB)
Non-trainable params: 0 (0.00 B)

```

Train the model

```

history = model.fit(
    x_train, y_train_cat,
    epochs=15,
    batch_size=64,
    validation_split=0.2,
    verbose=1
)

```

```

Epoch 1/15
750/750 8s 8ms/step - accuracy: 0.7704 - loss: 0.6630 - val
Epoch 2/15
750/750 5s 7ms/step - accuracy: 0.8598 - loss: 0.3872 - val
Epoch 3/15
750/750 6s 8ms/step - accuracy: 0.8756 - loss: 0.3365 - val
Epoch 4/15
750/750 5s 7ms/step - accuracy: 0.8860 - loss: 0.3113 - val
Epoch 5/15
750/750 6s 8ms/step - accuracy: 0.8922 - loss: 0.2873 - val
Epoch 6/15
750/750 5s 7ms/step - accuracy: 0.8987 - loss: 0.2724 - val
Epoch 7/15
750/750 5s 7ms/step - accuracy: 0.9063 - loss: 0.2519 - val
Epoch 8/15
750/750 6s 8ms/step - accuracy: 0.9031 - loss: 0.2529 - val
Epoch 9/15

```

```
750/750 ━━━━━━━━ 5s 7ms/step - accuracy: 0.9123 - loss: 0.2350 - val
Epoch 10/15
750/750 ━━━━━━ 6s 8ms/step - accuracy: 0.9158 - loss: 0.2240 - val
Epoch 11/15
750/750 ━━━━━━ 5s 7ms/step - accuracy: 0.9184 - loss: 0.2185 - val
Epoch 12/15
750/750 ━━━━━━ 7s 9ms/step - accuracy: 0.9198 - loss: 0.2054 - val
Epoch 13/15
750/750 ━━━━━━ 5s 7ms/step - accuracy: 0.9228 - loss: 0.2049 - val
Epoch 14/15
750/750 ━━━━━━ 6s 8ms/step - accuracy: 0.9269 - loss: 0.1909 - val
Epoch 15/15
750/750 ━━━━━━ 5s 7ms/step - accuracy: 0.9312 - loss: 0.1845 - val
```


Training Graphs

```
plt.figure(figsize=(14,5))

# Accuracy
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.title("Accuracy")
plt.legend()

# Loss
plt.subplot(1,2,2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title("Loss")
plt.legend()

plt.show()
```



Evaluate Test set

```
test_loss, test_acc = model.evaluate(x_test, y_test_cat, verbose=0)
print("Test Accuracy:", test_acc)
print("Test Loss:", test_loss)
```

```
Test Accuracy: 0.8891000151634216
Test Loss: 0.3550165891647339
```

Confusion Matrix

```
# Predictions
y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)

# Confusion matrix
cm = confusion_matrix(y_test, y_pred_classes)

plt.figure(figsize=(10,8))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=class_names,
            yticklabels=class_names)
```

```
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()

print(classification_report(y_test, y_pred_classes, target_names=class_names))
```


313/313

2s 5ms/step

Confusion Matrix

```
model.save("fashion_mnist_ffnn.h5")
```

```
WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `ke
```

Neural Network Image Classification Using Fashion-MNIST

Introduction This project demonstrates the application of a feedforward neural network (ANN) to an image classification task using the Fashion-MNIST dataset. Fashion-MNIST contains 70,000 grayscale images of clothing items across 10 categories. Compared to traditional machine learning, neural networks excel at extracting complex spatial features directly from pixel data, making them ideal for image classification problems.

This study focuses on building, training, and evaluating a simple neural network model while analyzing its performance using accuracy, loss, confusion matrix, and classification metrics.

Methodology

Dataset Preparation

- **Dataset:** Fashion-MNIST (70,000 images, 28×28 pixels, 10 classes)
 - **Split:** 60,000 training images, 10,000 test images
 - Sample images were visualized to understand pixel patterns and category variations.

T-shirt

Preprocessing

Pullover

- **Normalization:** Pixel values scaled from 0 to 255 to 0 to 1.
 - **One-hot encoding:** Converted labels into 10-class categorical vectors.
 - **Optional Augmentation:** Rotation, shifting transformations using

ImageSneaker

ImageDatabase

Bag

Model Architect

- Flatten Layer

- Dense^{macro avg}(256)

- weighted avg

- Dense(128, R)

• Dense(10, S)

- Dense(10, 50)

Model Compilation

- **Optimizer:** Adam
 - **Loss:** Categorical Crossentropy

-
- **Metrics:** Accuracy

Training

- 15 epochs
- Batch size: 64
- Validation split: 20%

Training and validation accuracy/loss were plotted for visual performance analysis.

Results

Training and Validation Curves The plots show:

- Increasing training accuracy
- Validation accuracy stabilizing around ~0.89
- Loss decreasing steadily

Test Performance

- **Test Accuracy:** ~0.889
- **Test Loss:** ~0.355

~~Confusion Matrix A confusion matrix was generated for deeper class-level insights.~~